

# RASPBERRY PI TUTORIAL FOR MACHINE LEARNING

SUBHRANSU S. BHATTACHARJEE

MAY 14, 2025




Australian  
National  
University

---

# DUE DILIGENCE:

## HAVE YOU SET UP THE PI?



Have you set up your Raspberry Pi's with the SD card?



Does the HDMI port work? Due to the lack of the monitors, we would like everyone to take a round-robin approach to using the Pi for each of the demos.



Let's get to the code repo and download the model.....

[https://github.com/1ssb/ANU\\_YOLO](https://github.com/1ssb/ANU_YOLO).


# TYPES OF LEARNING

---

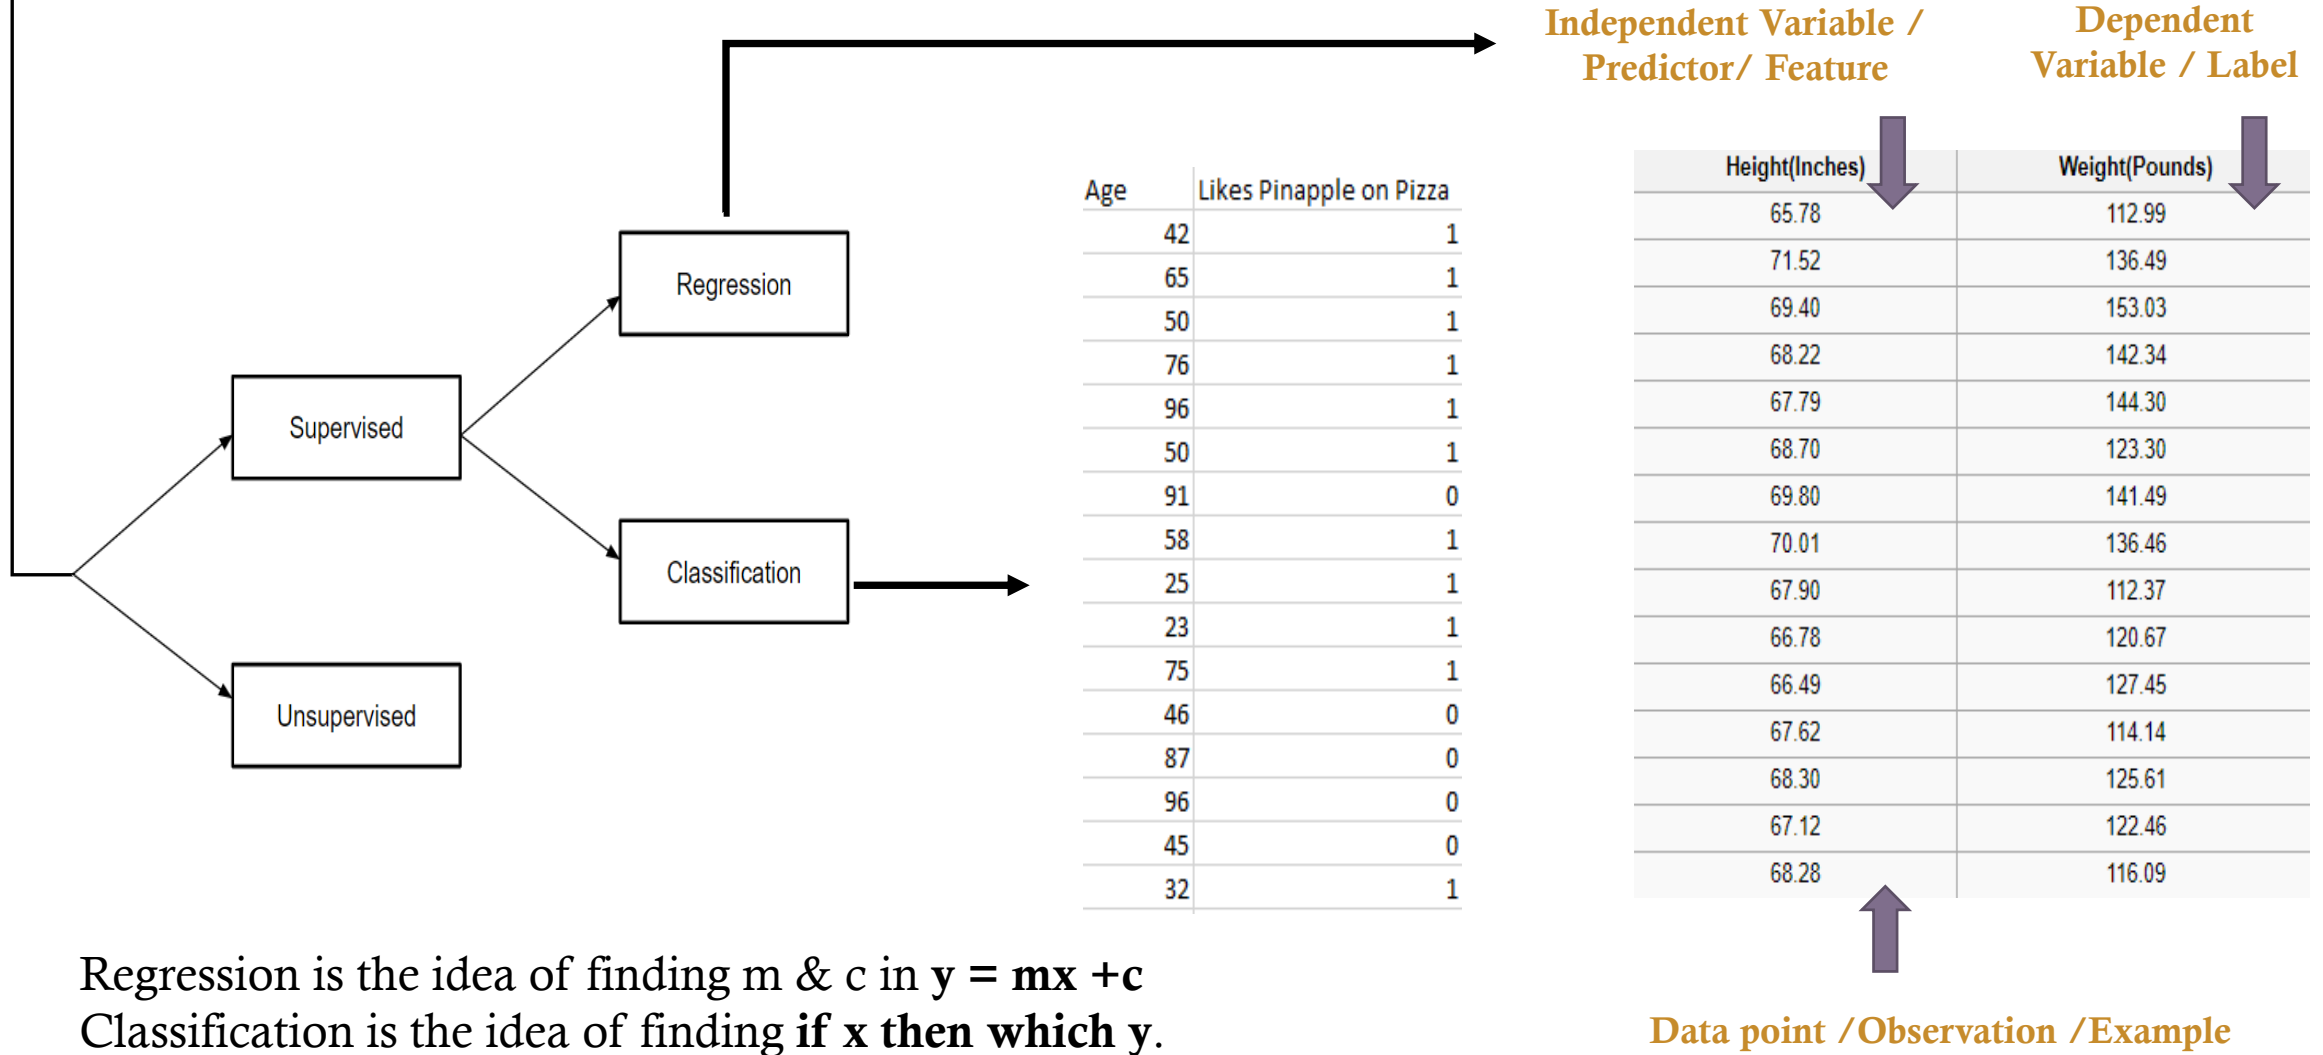
**Supervised learning** involves learning a function that maps an input to an output based on example input-output pair\*.



**Unsupervised learning** is used to draw inferences and find patterns from input data without references to labeled outcomes\*. Two main methods used in unsupervised learning include **clustering** and **dimensionality reduction**.



# TYPES OF LEARNING



# An Introduction to Learning using Computer Vision

---

## The purpose of Computer Vision

- **Automates perception:** Enables machines to “see” and interpret the world—key for robotics, autonomous vehicles, medical imaging, and analytics.
- **Scales human expertise:** Processes vast image/video data in real time, far beyond manual capabilities.
- **Drives innovation:** Powers AR/VR, smart cities, quality control in industry, and new human–machine interfaces.

## YOLO: Regression + Classification



1. **Regression:** Predicts precise bounding-box coordinates (x1, y1, x2, y2) directly from each grid cell.
  2. **Classification:** Simultaneously outputs class probabilities for each box (what object is inside it) for the class.
  3. **Unified model:** One network does both tasks in a single forward pass—fast and end-to-end trainable.
-

---

# YOLO: YOU ONLY LIVE ONCE

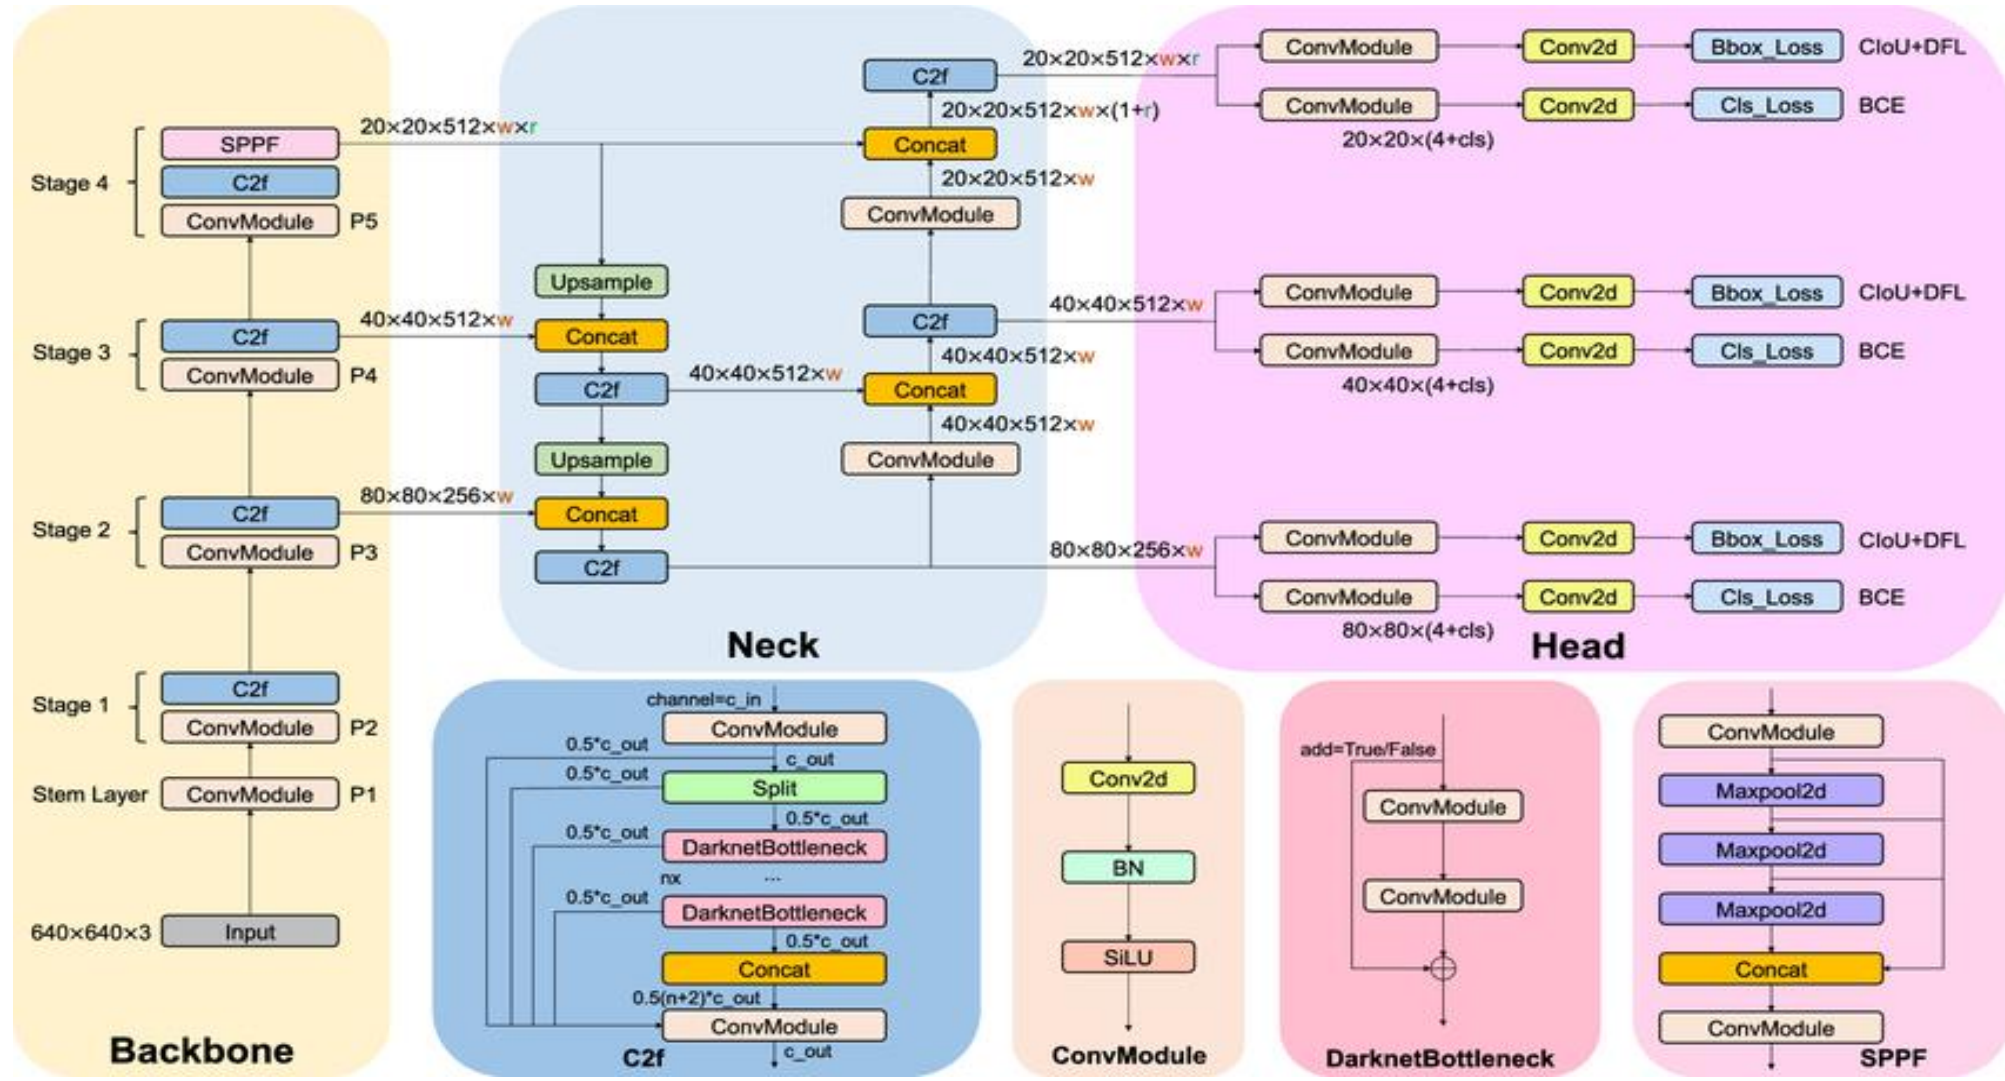
YOLO is a versatile algorithm trained on large-scale datasets to detect and classify objects in real time.

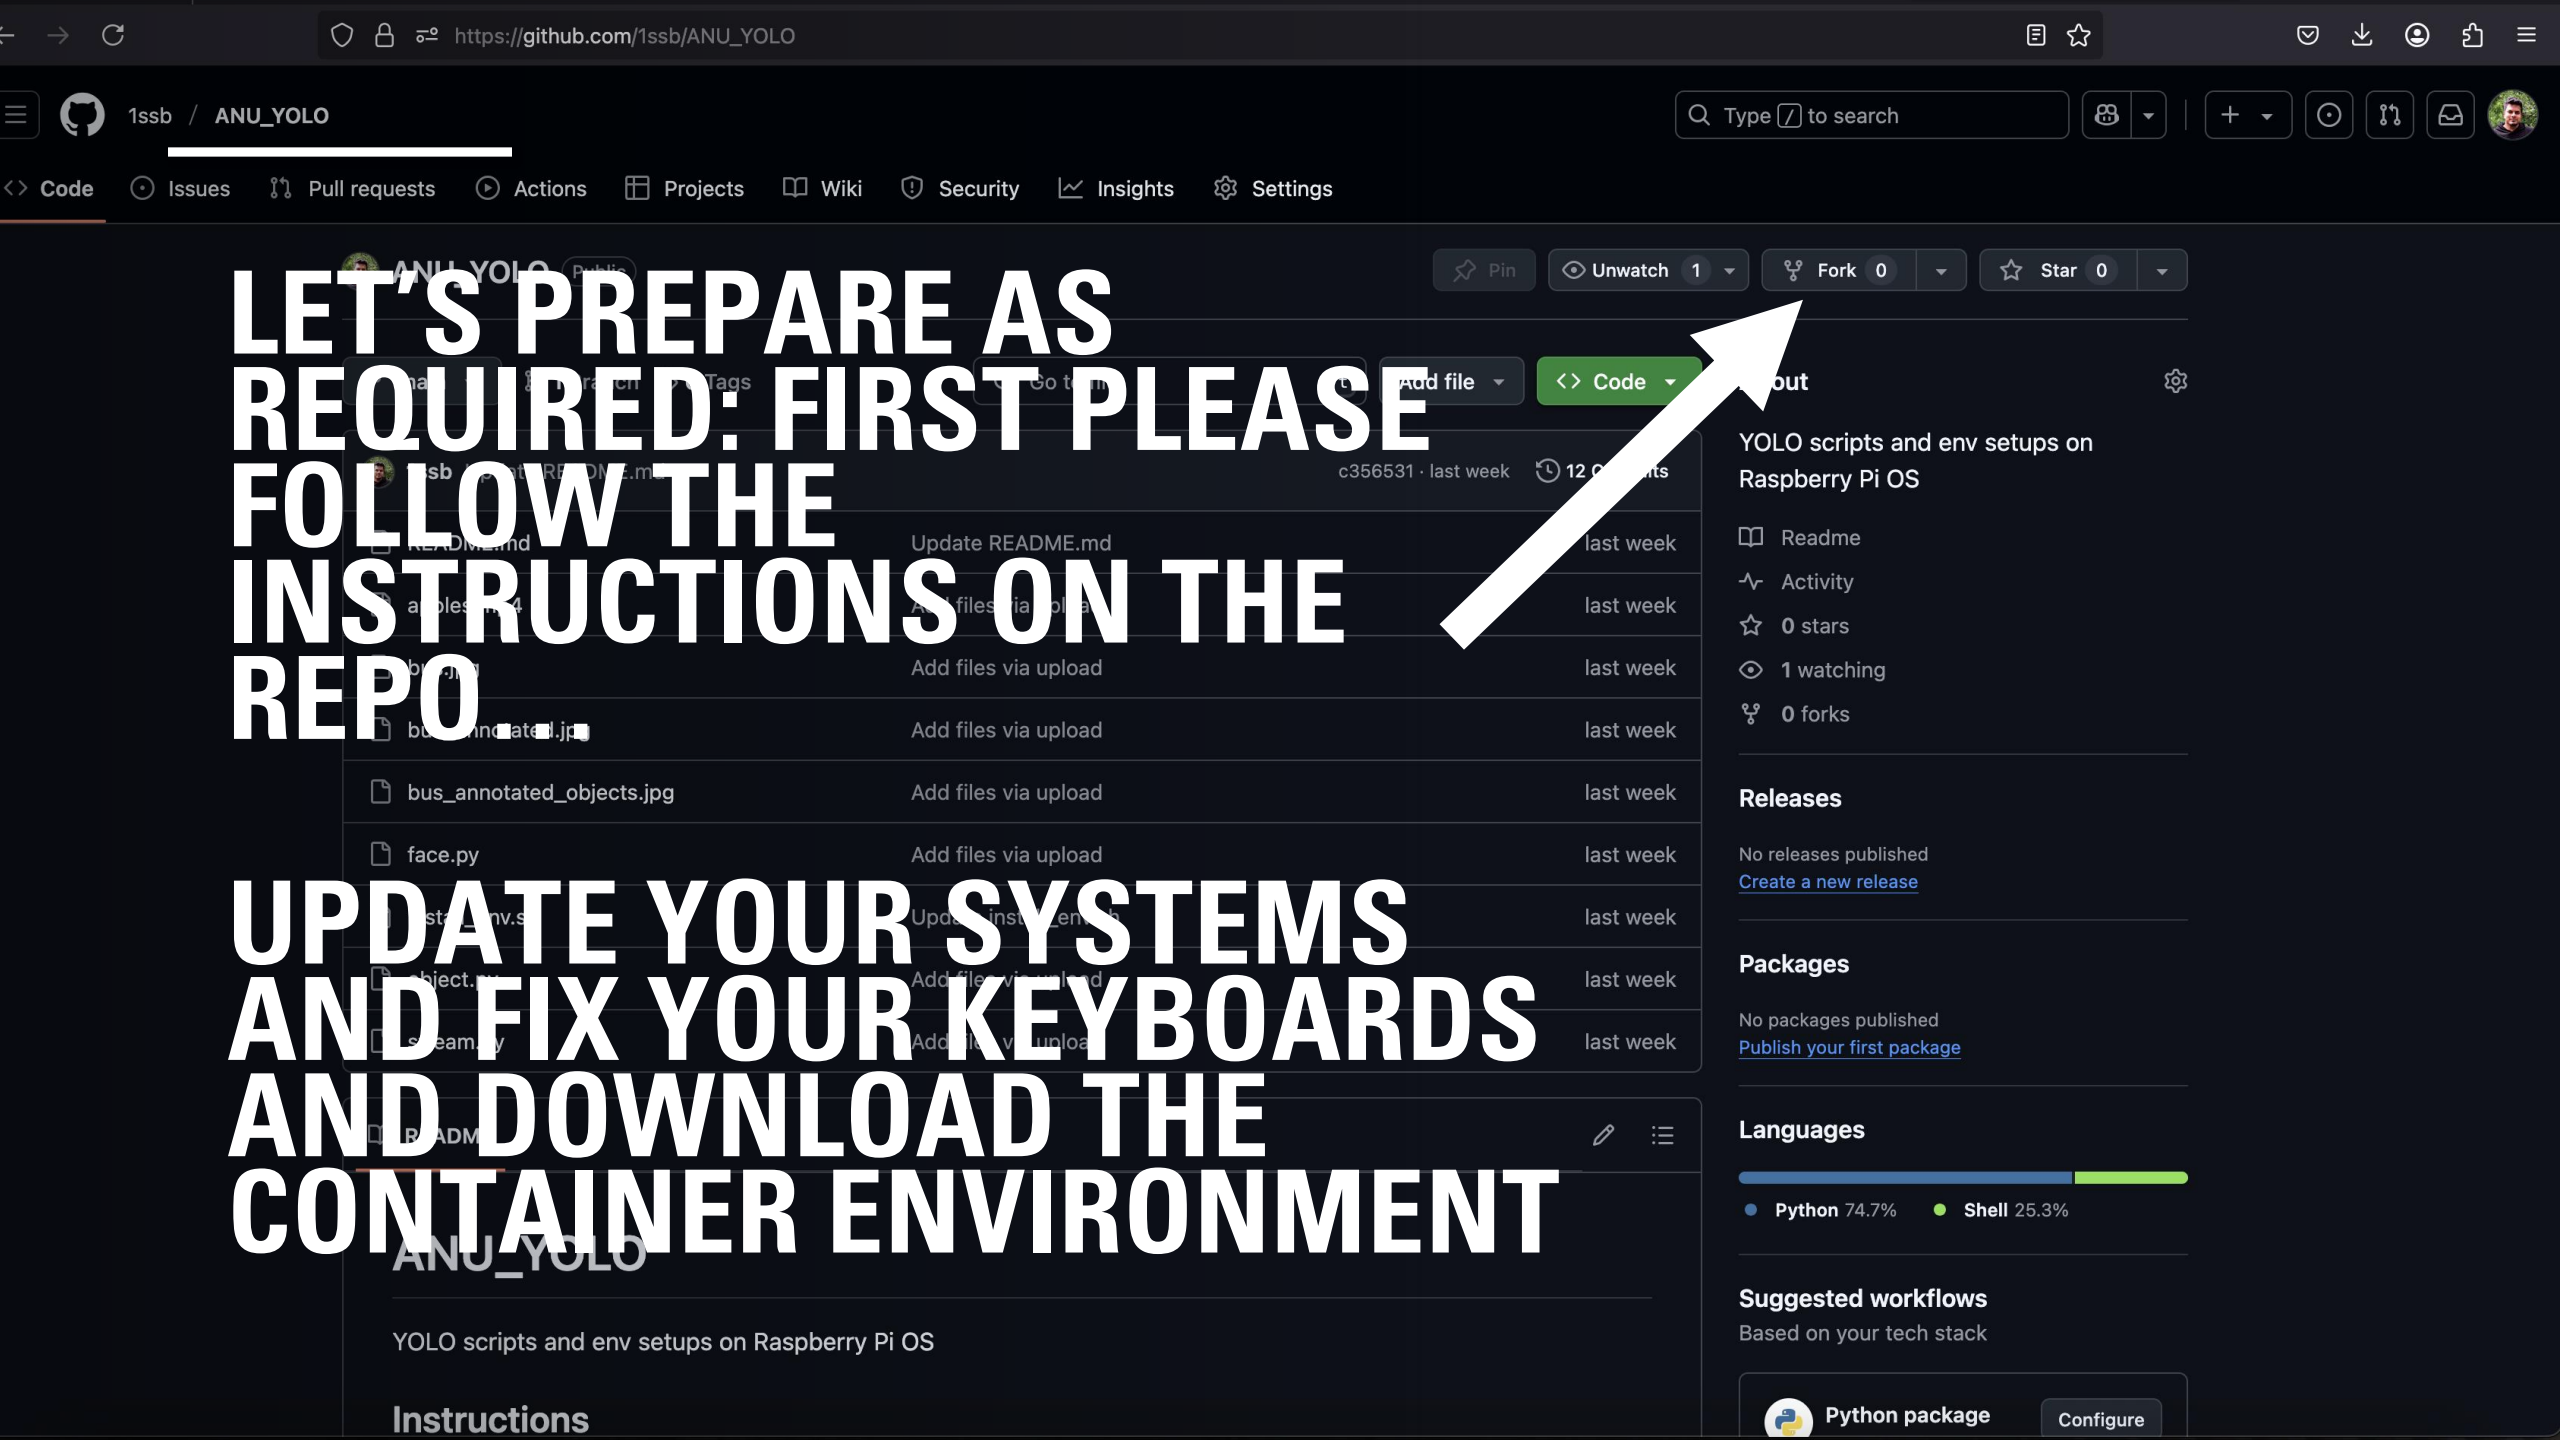
- **Single-pass detection:** Processes the entire image in one forward pass, unifying localization and classification without separate region proposals.
- **Grid-based outputs:** Splits the image into an  $S \times S$  grid; each cell predicts  $B$  bounding boxes ( $x, y, w, h$ ), a confidence score, and  $C$  class probabilities.
- **Real-time performance:** Runs at 30/60+ FPS on modern GPUs—and even on edge devices with the nano/small variants—while maintaining high accuracy and hence our exercise.





# YOLOv8 algorithm: Significantly more complex than how it started





LET'S PREPARE AS  
REQUIRED: FIRST PLEASE  
FOLLOW THE  
INSTRUCTIONS ON THE  
REPO...

UPDATE YOUR SYSTEMS  
AND FIX YOUR KEYBOARDS  
AND DOWNLOAD THE  
CONTAINER ENVIRONMENT

Type to search

Fork 0

Star 0

YOLO scripts and env setups on  
Raspberry Pi OS

Readme

Activity

0 stars

1 watching

0 forks

### Releases

No releases published  
[Create a new release](#)

### Packages

No packages published  
[Publish your first package](#)

### Languages

Python 74.7% Shell 25.3%

### Suggested workflows

Based on your tech stack



Python package


Configure

Instructions



---

# DEM01: OBJECT DETECTION



**Goal:** Identify and localize objects with bounding boxes and class labels.

**Example (YOLOv8n or Nano):**

*A street scene with “car”, “person”, and “bike” boxes labelled with confidence scores.*

**Key Steps:**

1. Input full image.
2. Split into grid.
3. Predict B boxes + confidence + class scores per cell. Maybe threshold the confidence.

---

BONUS: Will demonstrate the webcam object detection using stream.py if we have time



## DEMO 2: DETECT FACES

**Goal:** Locate and identify human faces within an image, returning bounding boxes and confidence scores.

### Key Steps:

- 1. Preprocessing:** Resize and normalize the input image (e.g., to  $416 \times 416$ ).
- 2. Feature Extraction:** Use a lightweight CNN backbone (e.g., Tiny-YOLO, MTCNN) to extract facial features.
- 3. Confidence Thresholding and NMS suppression:** Necessary in the process of calibration.

# GROUP TASK CHALLENGE: DETECT THE APPLES

---

**Task: Count the apples in the video in total.**

## **Option A: Detection + Tracking (HARDER but HIGH ACCURACY)**

- **Detect** each frame with a fruit-specialized object detector.
- **Assign IDs** to detections as they move across frames (e.g. via a simple tracker).
- **Aggregate** unique IDs to get the total count—prevents double-counting moving fruits.

## **Option B: Heuristic Frame Sampling (EASIER but LOW ACCURACY)**

- **Sample** one or more representative frames (e.g. every 10th frame or a key frame).
  - **Detect** fruits in each sampled frame.
  - **Take the maximum** number of detections seen in any single frame as your total—ideal when fruits don't reappear multiple times.
  - **JUSTIFY YOUR EXACT CHOICE** and sub-system design choices as comments.....
  - Use the Ultralytics documentation, link in repo.
-



---

# THANK YOU!

For your attention and time.

Questions?

**CONTACT ME**

W: [1ssb.github.io](https://1ssb.github.io)

E: [Subhransu.Bhattacharjee@anu.edu.au](mailto:Subhransu.Bhattacharjee@anu.edu.au)

